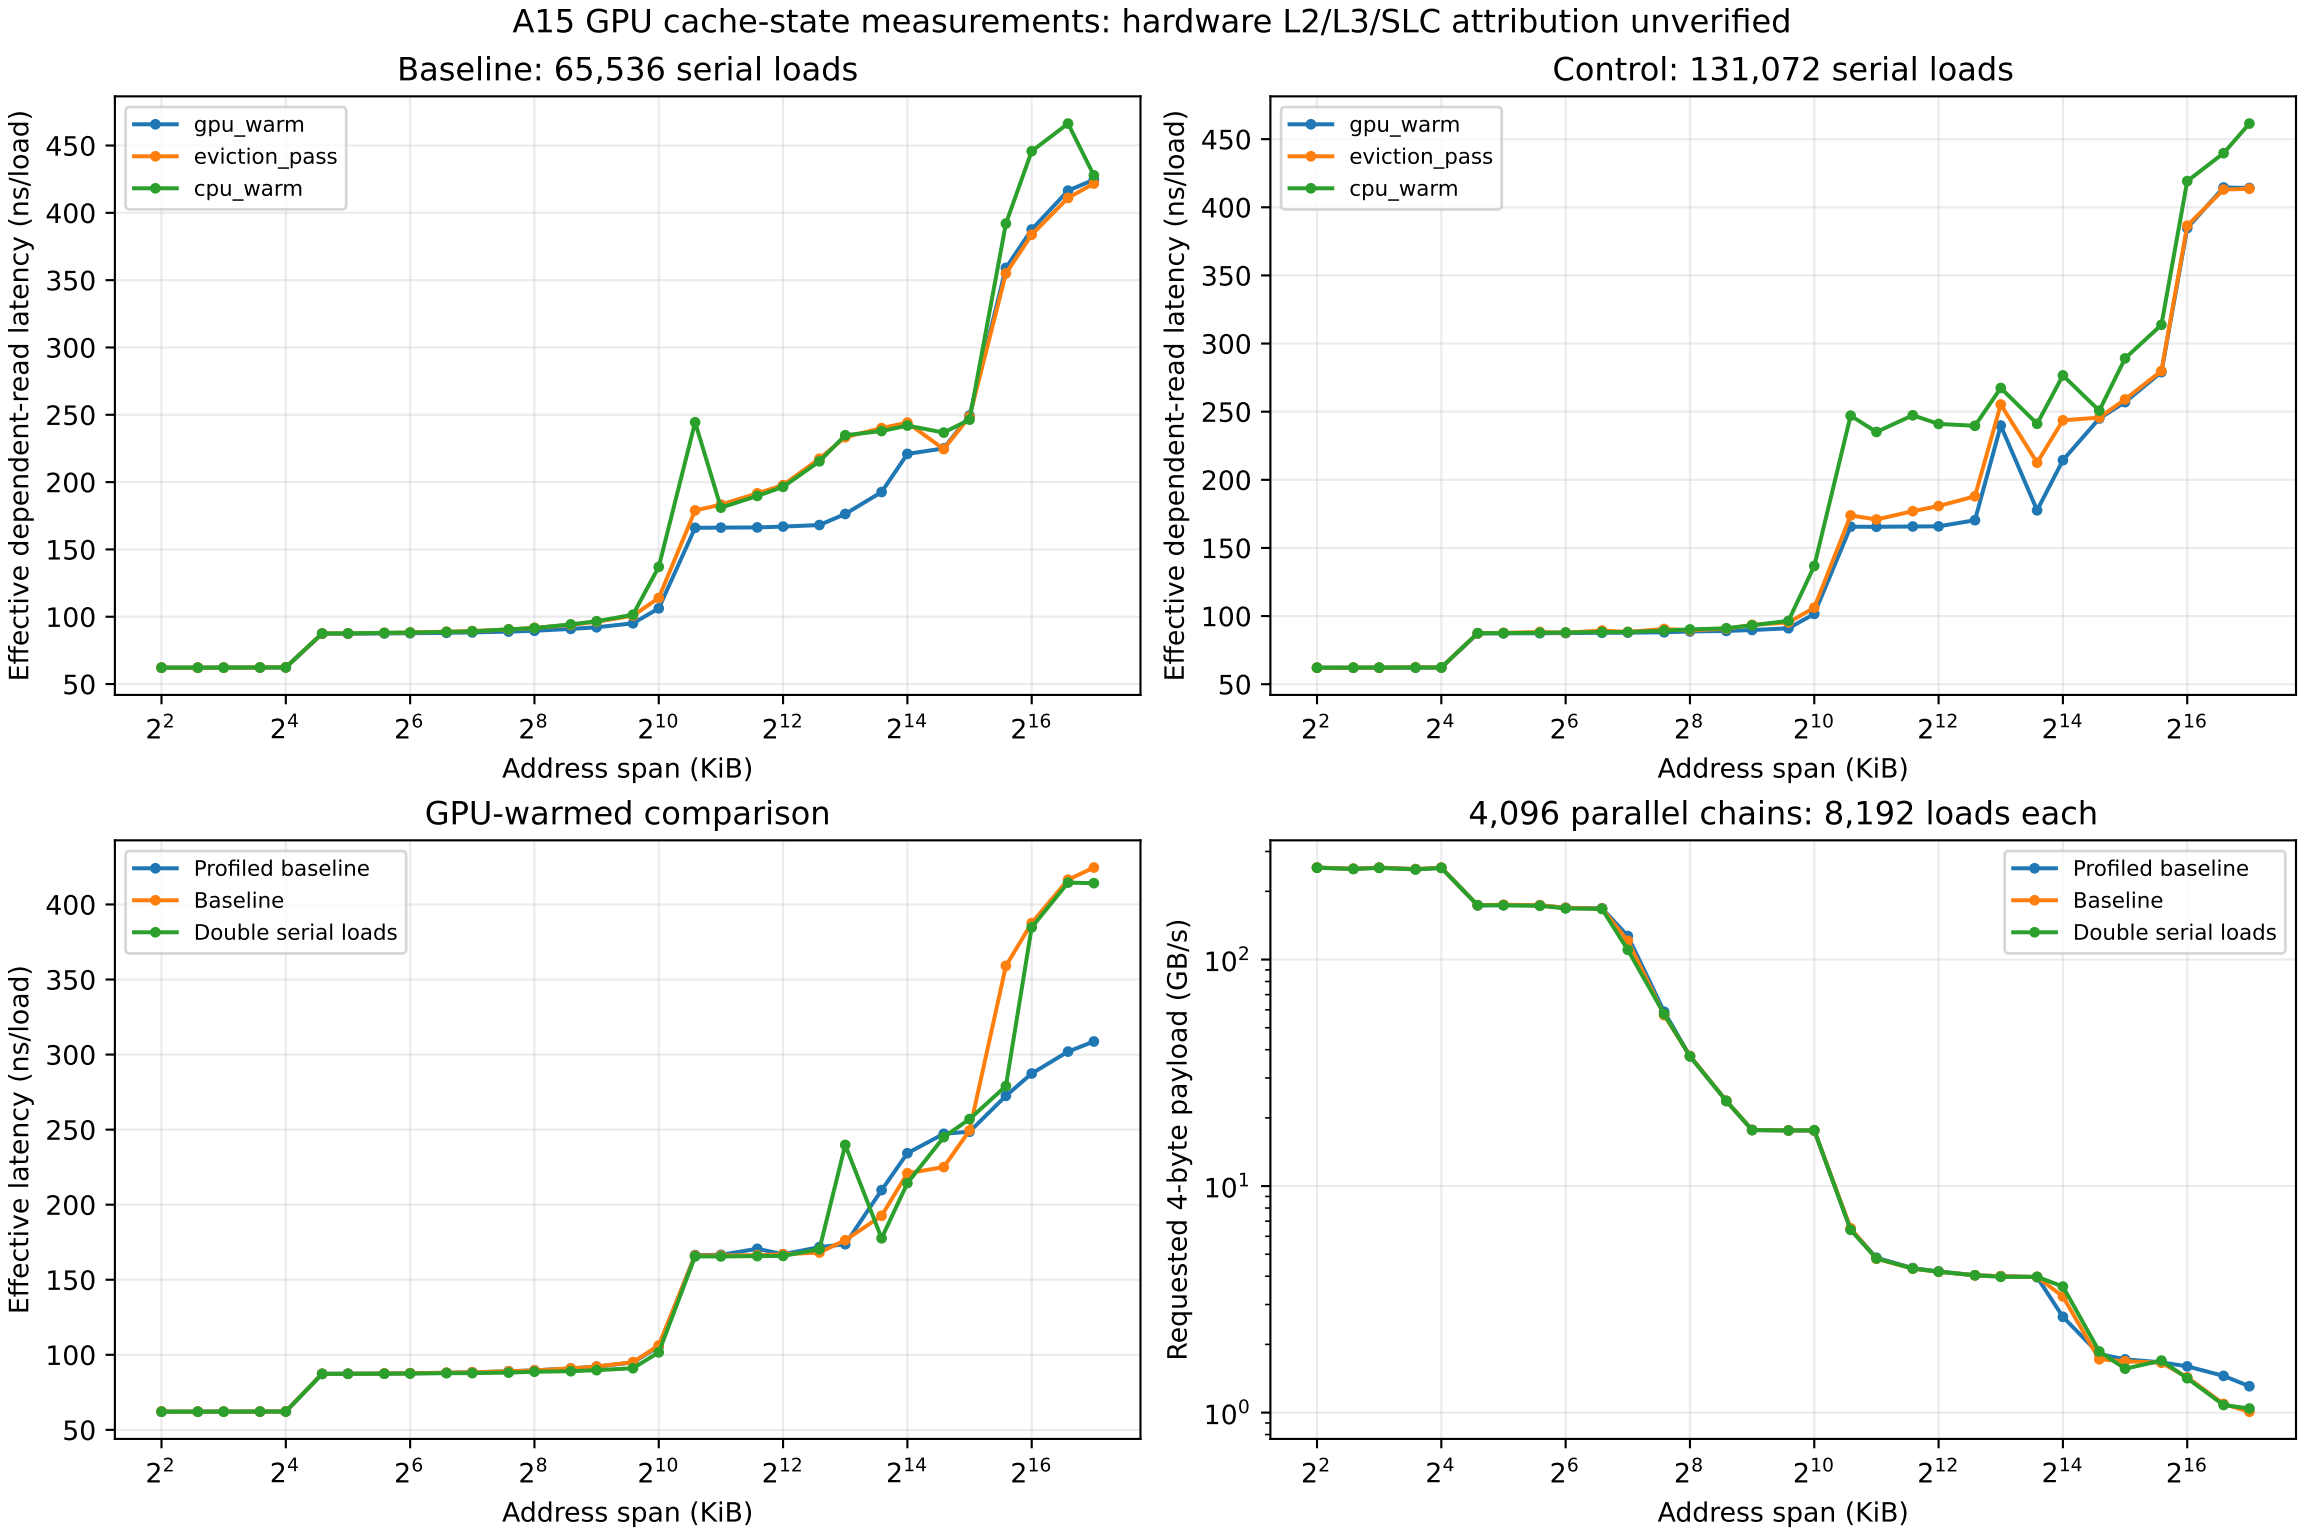
A15 GPU cache-state measurements: hardware L2/L3/SLC attribution unverified
Baseline: 65,536 serial loads
Control: 131,072 serial loads
GPU-warmed comparison
4,096 parallel chains: 8,192 loads each
gpu_warm
eviction_pass
cpu_warm
Profiled baseline
Baseline
Double serial loads
Effective dependent-read latency (ns/load)
Effective latency (ns/load)
Requested 4-byte payload (GB/s)
Address span (KiB)